

# 15-122: Principles of Imperative Computation, Fall 2017

## Written Homework 8

**Due:** Monday 23<sup>rd</sup> October, 2017 by 9pm

Name: Shaojie Zhang

Andrew ID: shaojiez

Section: O

This written homework covers amortized analysis, hash tables, and generics.

### Instructions

You can prepare your submission in one of two ways:

**Just edit (preferred)** Use any PDF editor (e.g., Preview on Mac, iAnnotate on mobile, Acrobat Pro installed on all non-CS cluster machines and most platforms) to typeset your answers in the given spaces — you can even draw pictures. *That's it.*

**Print and Scan** Alternatively, print this file, write your answers *neatly* by hand, and then scan it into a PDF file. *This is pretty labor-intensive.*

Once you have prepared your submission, submit it on Gradescope. You have unlimited submissions.

Question:	1	2	3	4	Total
Points:	2.5	4	1.5	4	12
Score:					

**Evaluation Summary** Once this homework is graded, you will be able to find a summary of your performance on Gradescope.

## 1. Amortized Analysis Revisited

Consider a special binary counter represented as  $k$  bits:  $b_{k-1}b_{k-2}\dots b_1b_0$ . For this special counter, the cost of flipping the  $i^{\text{th}}$  bit is  $2^i$  tokens. For example,  $b_0$  costs 1 token to flip,  $b_1$  costs 2 tokens to flip,  $b_2$  costs 4 tokens to flip, etc. We wish to analyze the cost of performing  $n = 2^k$  increments of this  $k$ -bit counter. (Note that  $k$  is *not* a constant.)

Observe that if we begin with our  $k$ -bit counter containing all 0s, and we increment  $n$  times, where  $n = 2^k$ , the final value stored in the counter will again be 0.

1pt

- 1.1** The worst case for a single increment of the counter is when every bit is set to 1. The increment then causes every bit to flip, the cost of which is

$$1 + 2 + 2^2 + 2^3 + \dots + 2^{k-1}$$

Explain in one or two sentences why this cost is  $O(n)$  — again take  $n = 2^k$ . (Hint: Find a closed form for the formula above.)

The above formula equals to  $2^k - 1$ . So  $O(2^k - 1)$  so it's  $O(n)$ .

1.5pts

- 1.2** Now, we will use amortized analysis to show that although the worst case for a single increment is  $O(n)$ , the amortized cost of a single increment is asymptotically less than this. Remember,  $n = 2^k$ .

Over the course of  $n$  increments, how many tokens in total does it cost to flip the  $i^{\text{th}}$  bit the necessary number of times?

**Solution:**  $n$

Based on your answer to the previous part, what is the total cost in tokens of performing  $n$  increments? (In other words, what is the total cost of flipping each of the  $k$  bits through  $n$  increments?) Write your answer as a function of  $n$  **only**. (Hint: what is  $k$  as a function of  $n$ ?)

**Solution:**  $n \log n$

Based on your answer above, what is the amortized cost of a single increment as a function of  $n$  **only**?

**Solution:**  $O(\log n)$  amortized

## 2. Hash Sets: Data Structure Invariants

A *hash set* is a hash table where there is no value associated with the keys: it is a convenient data structure to implement sets (of keys). The type `is_hset` defines hash sets similarly to separate-chaining hash tables. The code below checks that a given hash set is valid.

```
// typedef _____ key;
typedef struct chain_node chain;
struct chain_node {
    key data;
    chain* next;
};

struct hset_header {
    int size;           // number of elements stored in hash set
    int capacity;       // maximum number of chains in hash set
    chain*[] table;
};
typedef struct hset_header hset;

bool is_table_expected_length(chain*[] table, int length)
    //@assert \length(table) == length;
{ return true; }

bool is_hset(hset* H) {
    return H != NULL && H->capacity > 0 && H->size >= 0
        && is_table_expected_length(H->table, H->capacity);
}
```

An obvious data structure invariant of our hash set is that every element of a chain hashes to the index of that chain. Then, the above specification function is incomplete: we never test that the contents of the hash table satisfy this additional invariant. That is, we test only on the struct `hset`, and not on the properties of the array within.

On the next page, extend `is_hset` from above, adding a helper function to check that every element in the hash table belongs in the chain it is located in, and that each chain is non-cyclic. You should assume we will use the following two functions for hashing keys and for comparing them for equality:

```
int key_hash(key x);
bool key_equal(key x, key y);
```

Additionally, the function

```
int hash_index(hset* H, key k)
    /*@requires H != NULL && H->capacity > 0; @*/
    /*@ensures 0 <= \result && \result < H->capacity; @*/ ;
```

maps a hash to a valid index. It is provided for your convenience.

2pts

2.1 Note: your answer needs only to work for hash tables containing a few hundred million elements — do not worry about the number of elements exceeding `int_max()`.

```

bool has_valid_chains(hset* H)
// Preconditions (H != NULL, H->size >= 0, ...) omitted for space
{
    int nodecount = 0;

    for (int i = 0; i < H->capacity; i++) {
        // set p to the first node of chain i in table, if any

        chain* p = (H->table)[i];

        while (p != NULL) {
            key x = p->data;

            if (hash_index(H, x) != i)
                return false;

            nodecount++;

            if (nodecount > H->size)
                return false;

            p = p->next;
        }
    }

    if (nodecount != H->size)
        return false;

    return true;
}

bool is_hset(hset* H) {
    return H != NULL && H->capacity > 0 && H->size >= 0
        && is_table_expected_length(H->table, H->capacity)
        && has_valid_chains(H);
}

```

0.5pts

- 2.2 We generally don't care about the cost of specification functions, but what is the worst case complexity of `has_valid_chains` as a function of the number  $n$  of elements in the hash set?

Cost:  $O(\text{_____} n \text{_____})$

1.5pts

- 2.3 The updated function `is_hset` still falls short of flagging all possible invalid hash sets: nothing prevents a chain from containing multiple occurrences of an element. Given the above declarations, describe how you could check whether the hash set contains duplicate elements without allocating any extra memory. What is the cost?

use nested for loop for all the elements to check if they are the same,

i.e check the rest of the elements if there's duplicate

Cost:  $O(\text{_____} n^2 \text{_____})$

If you had a comparison function over keys, `int key_compare(key k1, key k2)`, how could you modify the behavior of the hash set to make the cost of finding duplicates asymptotically faster?

Change \_\_\_\_\_ the hash set to a sorted one

and then the compare only need to run through it once, so linear

Cost:  $O(\text{_____} n \text{_____})$

### 3. Hash Tables: Mapping Hash Values to Hash Table Indices

In our `hset` implementation, we require a library helper function `hash_index` that takes an element, computes its hash value using the client's `key_hash` function and converts this hash value to an integer. The first two functions below try to implement `hash_index` but have issues.

**0.5pts**

- 3.1** The following function has a bug. For one specific hash value `h`, this function does not return an index that is valid for a hash table. Identify the specific hash value.

```
int hash_index(hset* H, key k)
//@requires H != NULL && H->capacity > 0;
//@requires k != NULL;
//@ensures 0 <= \result && \result < H->capacity;
{
    int h = key_hash(H, k);
    return abs(h) % H->capacity;
}
```

**Solution:** This function fails when `h` = `int_min()`

**0.5pts**

- 3.2** The following function has an undesirable feature, although it always returns a valid index. Identify the flaw and, in one sentence, explain why it's a problem.

```
int hash_index(hset* H, key k)
//@requires H != NULL && H->capacity > 0;
//@requires k != NULL;
//@ensures 0 <= \result && \result < H->capacity;
{
    int h = key_hash(H, k);
    return h < 0 ? 0 : h % H->capacity;
}
```

There are too many elements hashed to index 0 and so the cost is high

0.5pts

- 3.3 Complete the following function so it avoids the problems in the previous two implementations of `hash_index`.

```
int hash_index(hset* H, key k)
//@requires H != NULL && H->capacity > 0;
//@requires k != NULL;
//@ensures 0 <= \result && \result < H->capacity;
{
    int h = key_hash(H, k);

    return (h < 0 ? _____abs(h+1)_____ : h) % H->capacity;
}
```

## 4. Generic Algorithms

A generic comparison function might be given a type as follows in C1:

```
typedef int compare_fn(void* x, void* y)
    //@ensures -1 <= \result && \result <= 1;
```

(Note: there's no precondition that **x** and **y** are necessarily non-NULL.)

If we're given such a function, we can treat **x** as being less than **y** if the function returns **-1**, treat **x** as being greater than **y** if the function returns **1**, and treat the two arguments as being equal if the function returns **0**.

Given such a comparison function, we can write a function to check that an array is sorted even though we don't know the type of its elements (as long as it is a pointer type):

```
bool is_sorted(void*[] A, int lo, int hi, compare_fn* comp)
    //@requires 0 <= lo && lo <= hi && hi <= \length(A) && comp != NULL;
```

1pt

4.1 Complete the generic binary search function below. You don't have access to generic variants of `lt_seg` and `gt_seg`. Remember that, for sorted integer arrays, `gt_seg(x, A, 0, lo)` was equivalent to `lo == 0 || A[lo - 1] < x`.

```
int binsearch_generic(void* x, void*[] A, int n, compare_fn* cmp)
    //@requires 0 <= n && n <= \length(A) && cmp != NULL;
    //@requires is_sorted(A, 0, n, cmp);
{
    int lo = 0;
    int hi = n;

    while (lo < hi)
        //@loop_invariant 0 <= lo && lo <= hi && hi <= n;

        //@loop_invariant lo == 0 || (*cmp)(A[lo-1], x) == -1;

        //@loop_invariant hi == n || (*cmp)(A[hi], x) == 1;
    {
        int mid = lo + (hi - lo)/2;

        int c = (*cmp)(A[mid], x);

        if (c == 0) return mid;
        if (c < 0) lo = mid + 1;
        else hi = mid;
    }
    return -1;
}
```



Suppose you have a generic sorting function, with the following contract:

```
void sort_generic(void*[] A, int lo, int hi, compare_fn* cmp)
  //@requires 0 <= lo && lo <= hi && hi <= \length(A) && cmp != NULL;
  //@ensures is_sorted(A, lo, hi, cmp);
```

1pt

- 4.2 Write an integer comparison function `compare_ints` that can be used with this generic sorting function, which you should assume is already written. You can leave out the postcondition that the result of `compare_ints` is between `-1` and `1` inclusive. However, the contracts on your `compare_ints` function *must* be sufficient to ensure that no precondition-passing call to `compare_ints` can possibly cause a memory error.

```
int compare_ints(void* x, void* y)

//@requires x != NULL && \hastag( __int*, x
                                     );
//@requires y != NULL && \hastag( __int*, y
                                     );
{
    if ( __*(int*) x < *(int*) y ) return __-1;
    if ( __*(int*) x > *(int*) y ) return __1;
    return __0;
}
```

2pts

- 4.3 Using the above generic sorting function and `compare_ints`, fill in the body of the `sort_ints` function below so that it will sort the array `A` of integers using the generic sort function specified above. You can omit loop invariants. But of course, when you call `sort_generic`, the preconditions of `compare_ints` must be satisfied by any two elements of the array `B`.

```
void sort_ints(int[] A, int n)
//@requires \length(A) == n;
{
    // Allocate a temporary generic array of the same size as A

    void*[] B = ____ alloc_array(void*, n) ____;

    // Store a copy of each element in A into B

    for (int i = 0; i < n; i++) {
        int * x = alloc(int);
        *x = A[i];
        B[i] = (void*) x;
    }

    // Sort B using sort_generic and compare_ints from part b
    sort_generic(B, 0, n, &compare_ints);

    // Copy the sorted ints in your generic array B into array A

    for (int i = 0; i < n; i++) {
        A[i] = *(int*) B[i];
    }

}
```